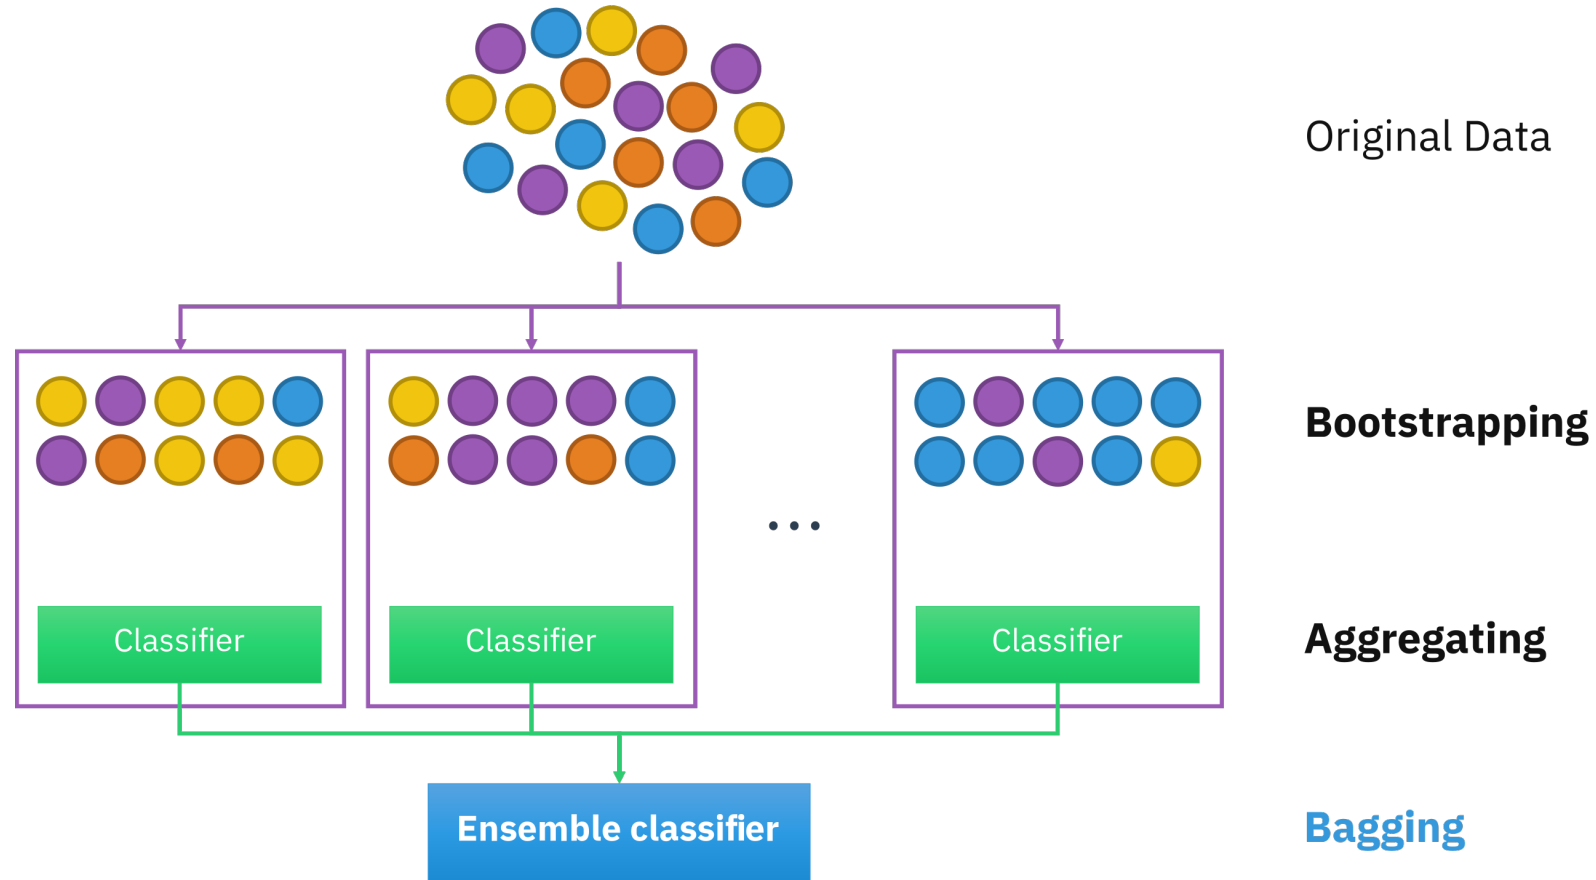


Day 4

# From Trees to Forests

# Bagging: Random Forests

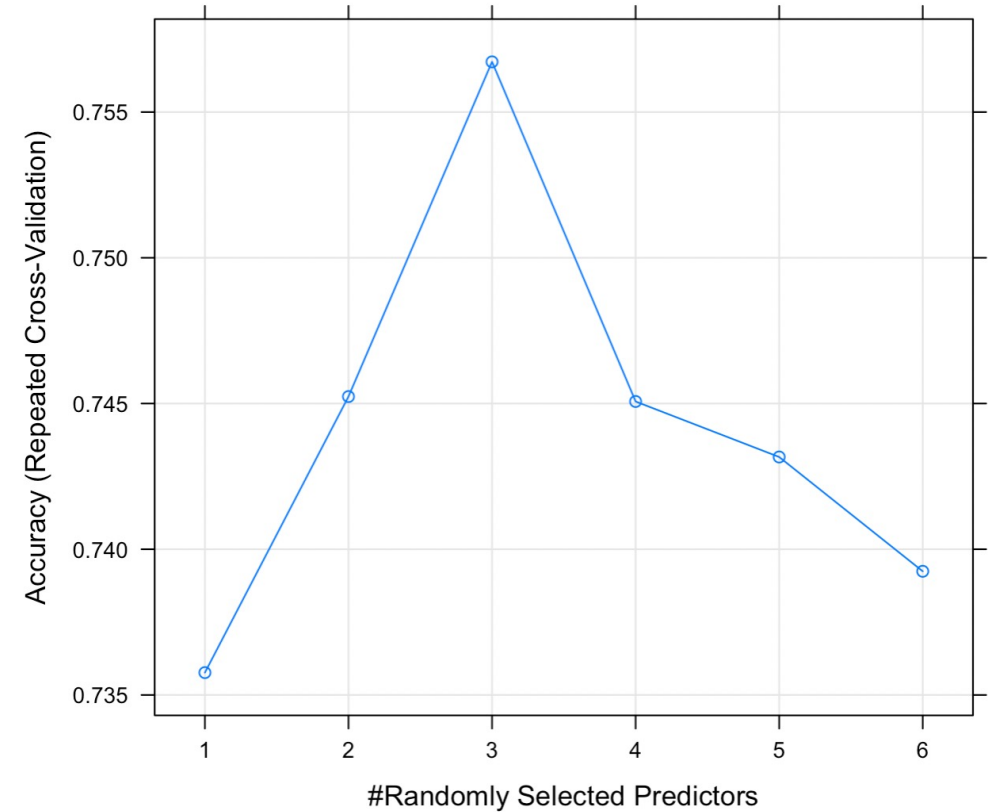


# Parameters of the model

- The ***mtry*** parameter controls how many of the input features a decision tree has available to consider at any given point in time.
- Since different sets of features will be available to different decision trees at different points, it will be (nearly) impossible for all of your trees to look exactly the same.
- Predictions are then obtained by taking the average over posterior probabilities or taking the “majority vote”

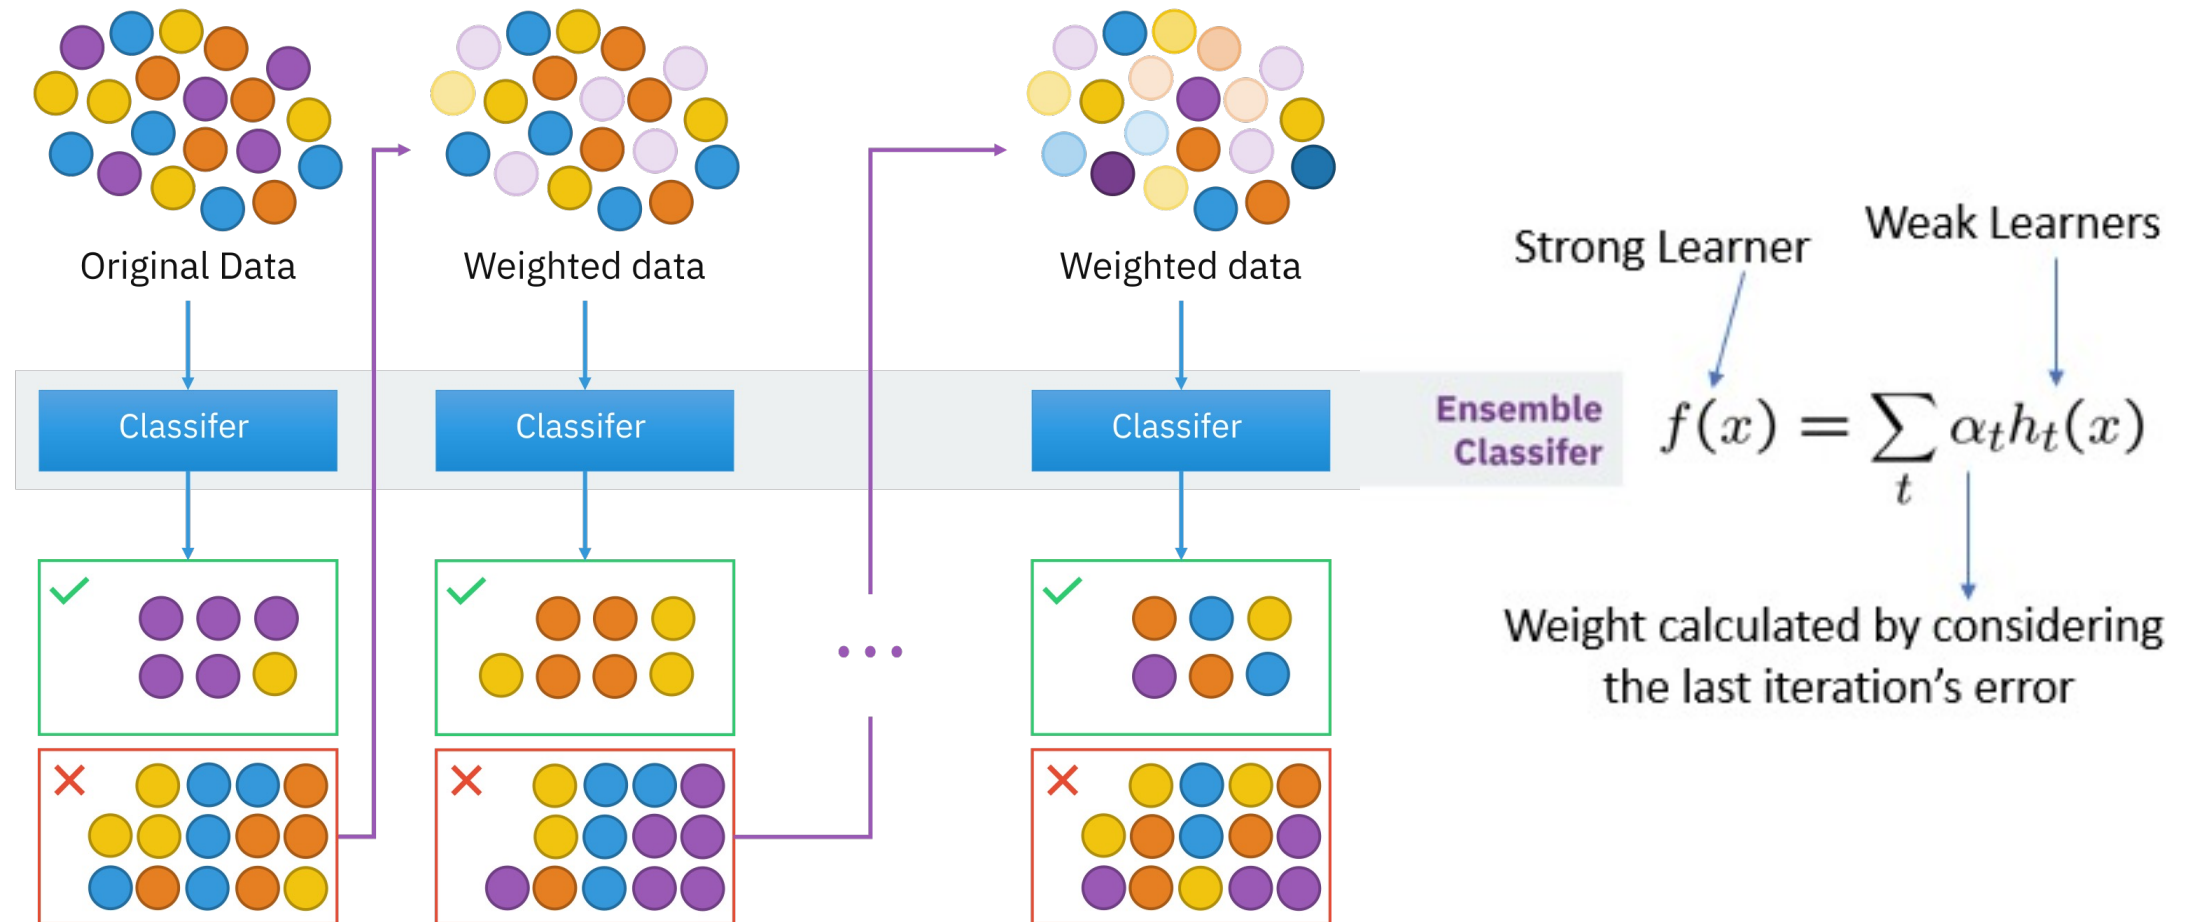
# Setting up a search grid

```
tunegrid <- expand.grid(  
  mtry = c(1:6)  
)  
  
trcontrol <- trainControl(method = "repeatedcv",  
  number = 5,  
  repeats = 1,  
  search = "grid")  
  
rf_train = caret::train(labels~ .,  
  data=pp_train,  
  trControl = trcontrol,  
  tuneGrid = tunegrid,  
  method = 'parRF')
```

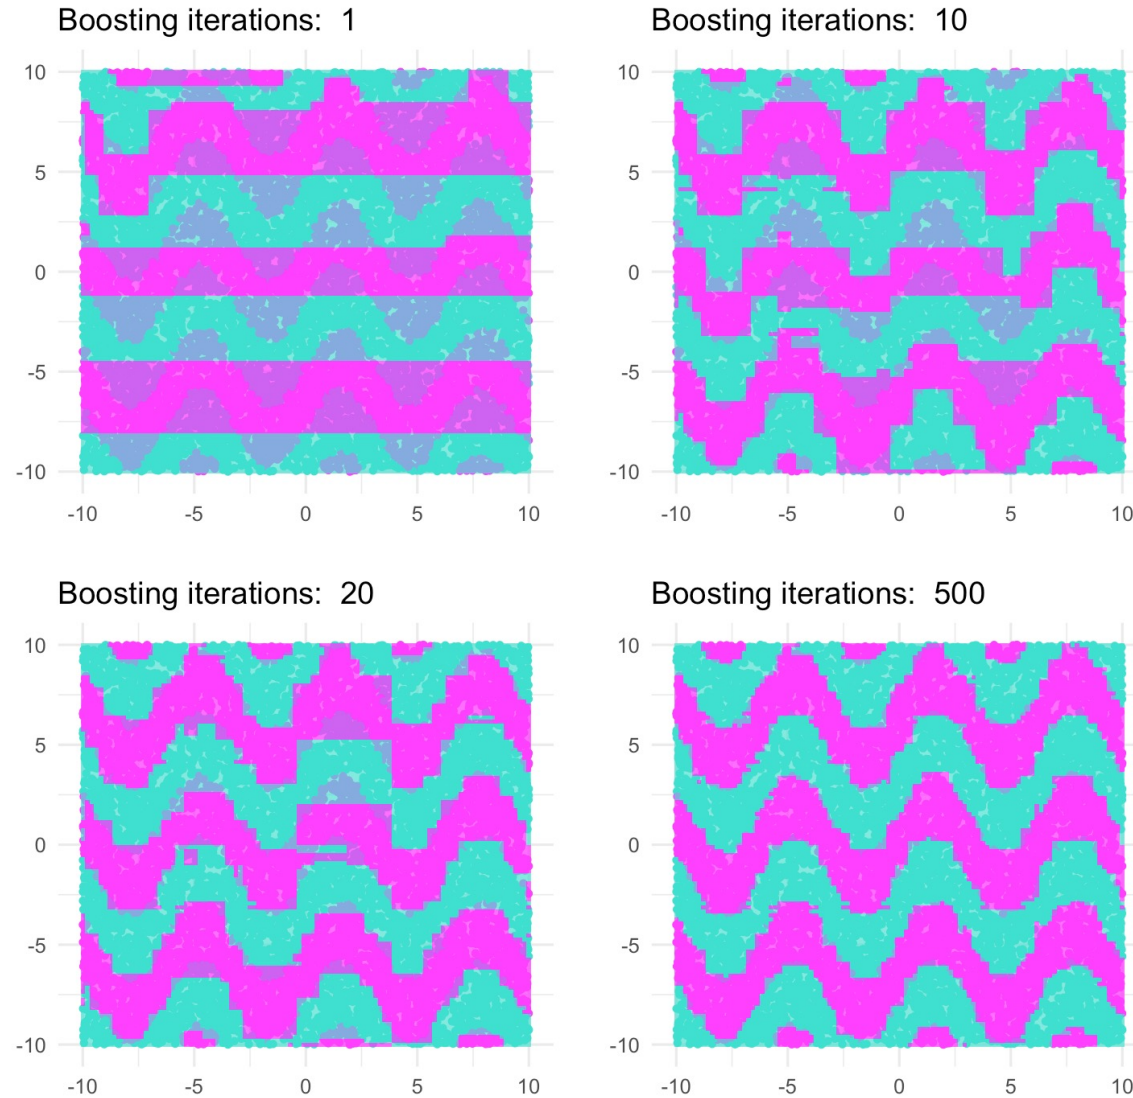


Boosted trees

# Boosting



# Number of boosting iterations





# XGBoost - tunegrid

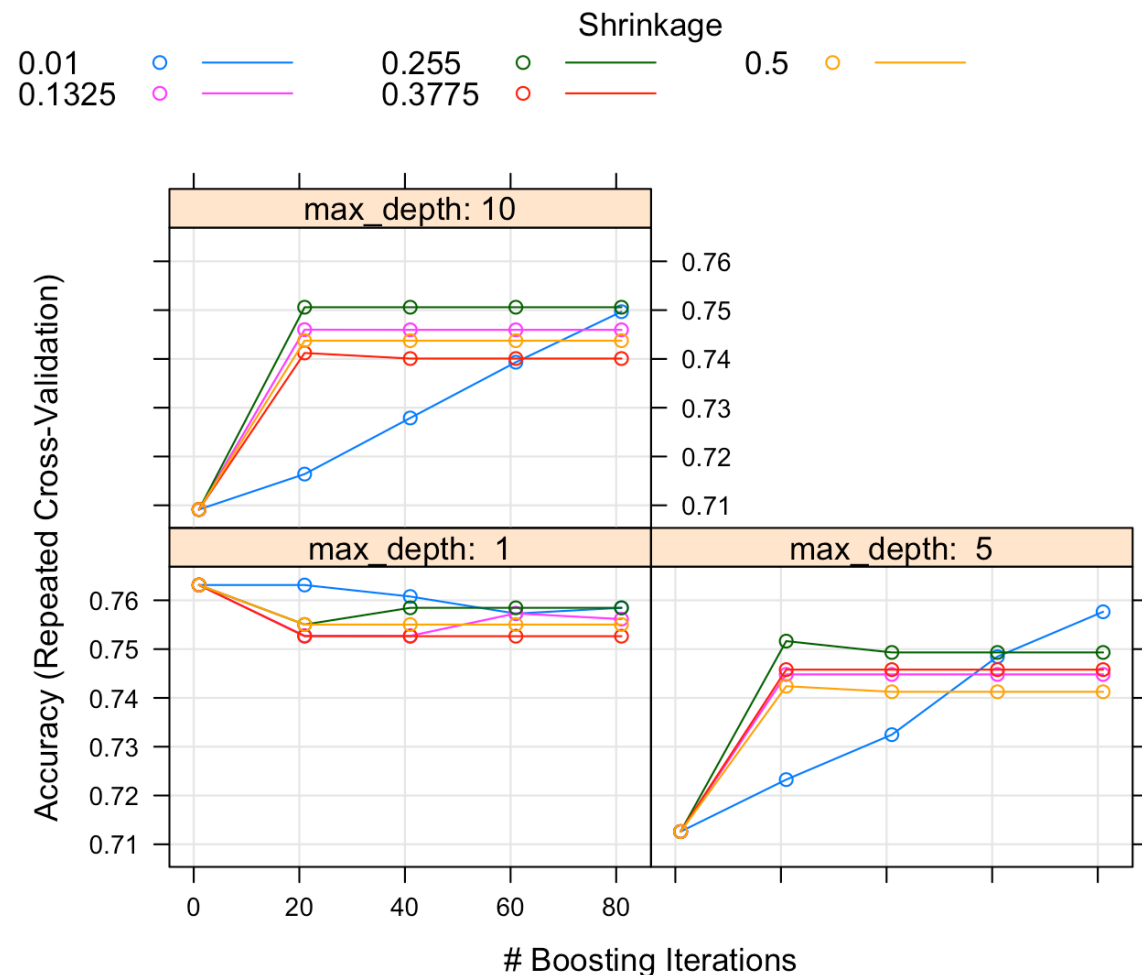
```
tunegrid <- expand.grid(nrounds = seq(1,100,by=20),  
                        max_depth = c(1,5,10),  
                        eta = seq(0.01,0.5,length.out = 5),  
                        gamma = 1,  
                        colsample_bytree = 1,  
                        min_child_weight = 1,  
                        subsample = 1)
```

```
trcontrol <- trainControl(method = "repeatedcv",  
                           number = 5,  
                           repeats = 5,  
                           search = "random",  
                           savePredictions = T)
```

# XGBoost – hyperparameter tuning


```
xgb_train = train(labels~ .,  
                  data=pp_train,  
                  trControl = trcontrol,  
                  tuneGrid = tunegrid,  
                  method = "xgbTree")
```

```
plot(xgb_train)
```



# Confusion Matrix

	XGBoost
Sensitivity	0.6400000
Specificity	0.8750000
Pos Pred Value	0.7272727
Neg Pred Value	0.8235294
Precision	0.7272727
Recall	0.6400000
F1	0.6808511
Prevalence	0.3424658
Detection Rate	0.2191781
Detection Prevalence	0.3013699
Balanced Accuracy	0.7575000



# Summary of Models

	Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision	Recall	F1	Prevalence	Detection Rate	Detection Prevalence	Balanced Accuracy
RPart	0.48	0.8750000	0.6666667	0.7636364	0.6666667	0.48	0.5581395	0.3424658	0.1643836	0.2465753	0.6775000
RandomForest	0.64	0.8333333	0.6666667	0.8163265	0.6666667	0.64	0.6530612	0.3424658	0.2191781	0.3287671	0.7366667
XGBoost	0.64	0.8750000	0.7272727	0.8235294	0.7272727	0.64	0.6808511	0.3424658	0.2191781	0.3013699	0.7575000

- Sensitivity or recall** (true positive rate) is the probability of a positive test result, conditioned on the individual truly being positive.

- Specificity** (true negative rate) is the probability of a negative test result, conditioned on the individual truly being negative.

- The **positive and negative predictive values** are the proportions of positive and negative results that are true positive and true negative results, respectively. Positive predictive value is also called **Precision**.

- The **F1** score is defined as the harmonic mean of precision and recall.







- Detection rate** is the proportion of the whole sample where the events were detected correctly.

- Detection prevalence** is the proportion of the whole sample that were classified as the focal category (= “not surviving” in our case).

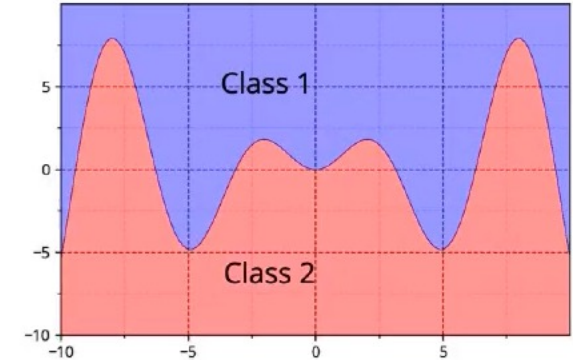


- Balanced Accuracy** is the average of specificity and sensitivity.

# Neural Networks

# Neural Networks: universal approximation theorem





# Neural Networks: universal approximation theorem





The *universal approximation theorem* states that any continuous function ( = classifier) can be approximated by a feed-forward neural network!

But: (heuristic) algorithms like back-propagation need to be employed to train the network.



# Network Architecture

- **Input Layer:** Receives the initial data.
- **Hidden Layers:** Intermediate layers performing computations/feature transformations.
- **Output Layer:** Produces the final prediction/output.





# Hyperparameters (list not exhaustive)

- **Learning Rate ( $\eta$ ):**
  - Controls the step size during the optimization process.
  - Affects how quickly the model converges to a minimum (too high can overshoot, too low can stall).
- **Regularization Parameters:**
  - **L2 Regularization (Ridge):** Penalizes large weights, encouraging simpler models.
  - **L1 Regularization (Lasso):** Can lead to sparse models by zeroing out weights.
- **Number of Layers:**
  - **Depth of Network:** The number of hidden layers in the network.
  - More layers can capture more complex patterns but increase computational complexity and training time.
- **Number of Neurons per Layer:**
  - Determines the capacity of each hidden layer.
  - More neurons can model more complex functions but may increase risk of overfitting.
- **Activation Functions:**
  - **ReLU (Rectified Linear Unit):** Commonly used, helps mitigate the vanishing gradient problem.
  - **Sigmoid:** Squashes input between 0 and 1, useful for binary classification.
  - **Tanh:** Squashes input between -1 and 1, zero-centered.

# Nnet – hyperparameter tuning


```
train_control <- trainControl(  
  method = "cv",  
  number = 20,  
  classProbs = TRUE)
```

```
tune_grid <- expand.grid(  
  size = (3:8),  
  decay = c(0.01, 0.1, 0.2))
```

# Nnet – hyperparameter tuning

```
mod_nnet = caret::train(labels~ .,  
  data=pp_train,  
  method = "nnet",  
  trControl = train_control,  
  tuneGrid = tune_grid,  
  trace = FALSE)
```

```
plot(mod_nnet)
```

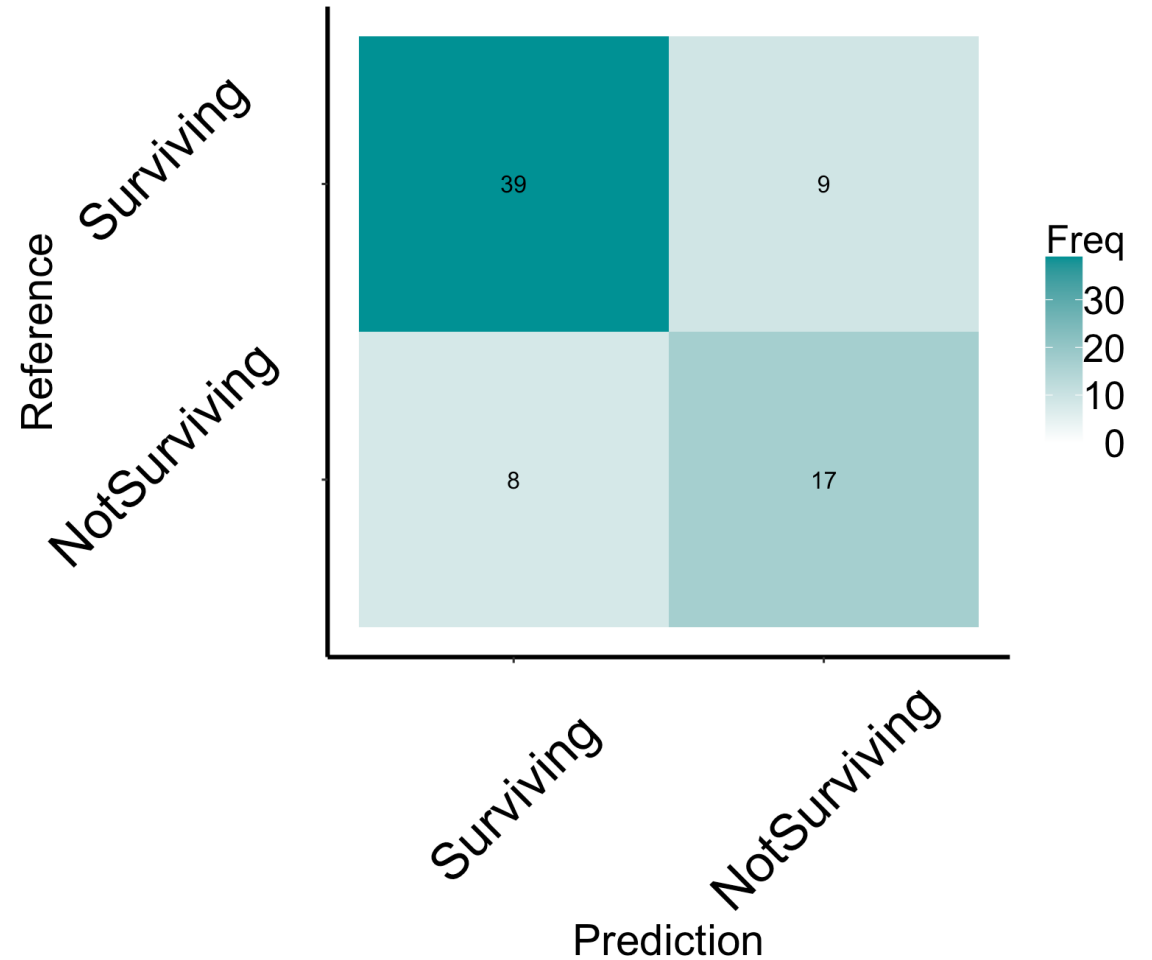


# Nnet – performance

```
pred <- predict(mod_nnet, pp_test)
(cm <- confusionMatrix(pred,
pp_test$labels))
```

	Reference	
Prediction	NotSurviving	Surviving
NotSurviving	17	9
Surviving	8	39

Accuracy : 0.7671  
95% CI : (0.6535, 0.8581)  
No Information Rate : 0.6575  
P-Value [Acc > NIR] : 0.02939



# Practicals 3.1

Trying different algorithms

# Balancing Data

- Data imbalance occurs when one class of the target variable is significantly more frequent than others.
- Techniques like oversampling, undersampling, SMOTE (Synthetic Minority Over-sampling Technique), and class weights can help address data imbalance.
- Balancing data enhances the model's ability to learn from minority class examples and improves overall performance.

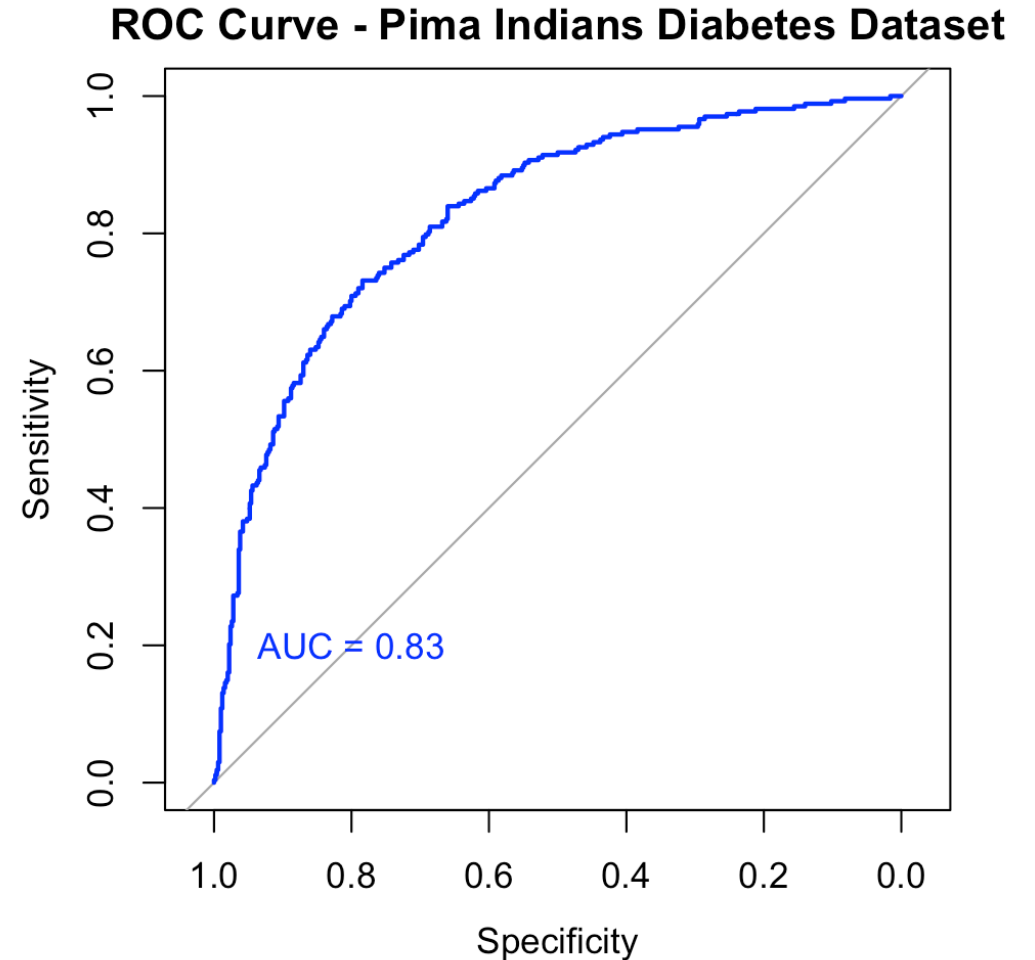
# Loss Functions

- A function that measures the model's prediction accuracy by comparing predicted values with actual values during training.
- Common loss functions include:
  - Residual Sum of Squares (RSS),
  - Mean Squared Error (MSE),
  - Cross-Entropy Loss,
  - L1 Loss (Mean Absolute Error)
  - ...
- Optimizing the loss function helps in improving the model's performance.

# ROC Curves

(Receiver Operating Characteristic)

- Graphical representation of the trade-off between true positive rate (sensitivity) and false positive rate (1 - specificity) for varying classification thresholds.
- Summarizes the model's performance across different threshold values.
- AUC (Area Under the Curve) is often used to quantify the ROC curve's performance.





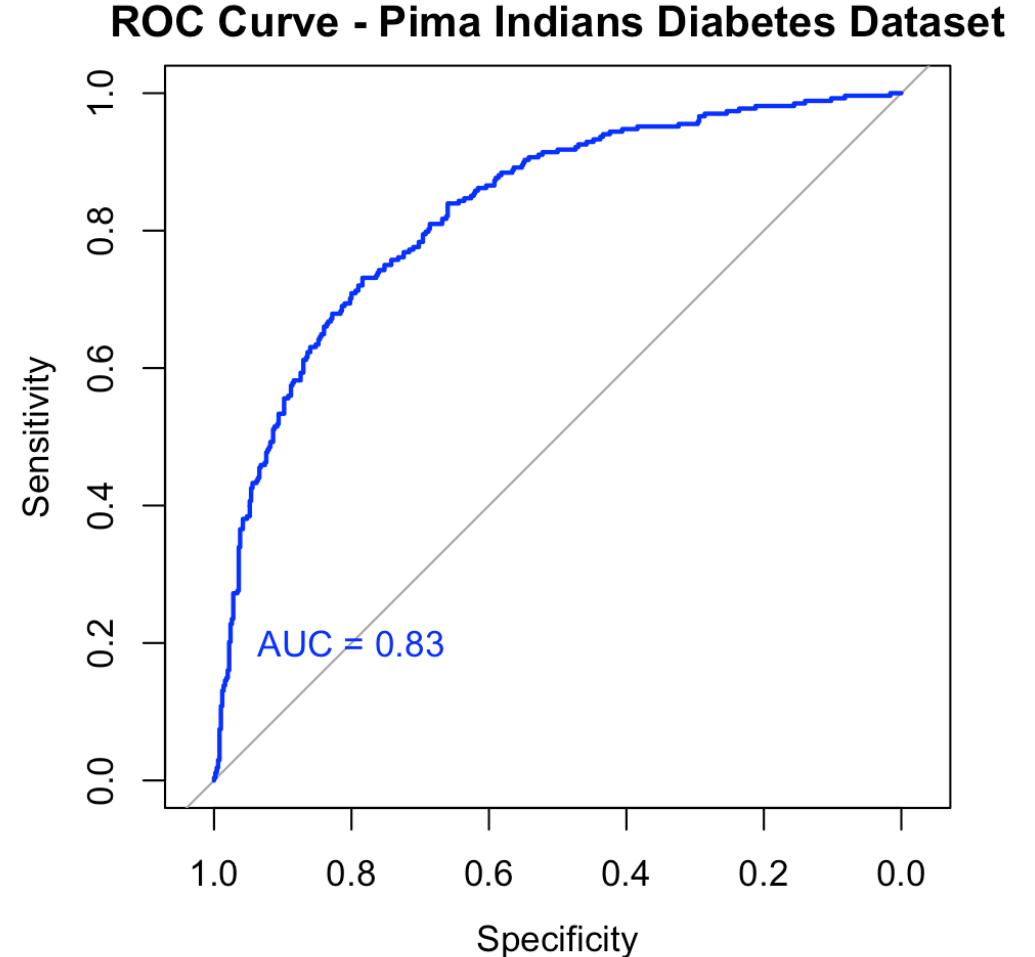
# ROC Curves

(Receiver Operating Characteristic)

```
# Make predictions on the training set
predictions <- predict(model,
                        PimaIndiansDiabetes,
                        type = "prob")

resp = ifelse(
  PimaIndiansDiabetes$diabetes == "pos",
  1, 0)

# Calculate ROC curve
roc_curve <- roc(response = resp,
                 predictor = predictions$pos)
```



# Supervised Learning

- Involves training a model on labeled data where the input features and their corresponding output labels are provided.
- The model learns to map inputs to outputs by finding patterns and relationships in the labeled data.
- Requires a clear objective, such as classification or regression, to minimize a specific loss function during training.
- Examples include linear regression, logistic regression, support vector machines, and neural networks.

# Unsupervised Learning

- Involves training a model on unlabeled data where only input features are provided without explicit output labels.
- The model learns to find hidden patterns, structures, or relationships in the data without specific guidance on what to look for.
- Often used for tasks like clustering, dimensionality reduction, and anomaly detection.
- Does not rely on a predefined objective or labeled data, allowing for more exploratory data analysis and pattern discovery.

## 1. Clustering:

1. K-means clustering
2. Hierarchical clustering

## 2. Dimensionality Reduction:

1. Principal Component Analysis (PCA)
2. t-Distributed Stochastic Neighbor Embedding (t-SNE)
3. UMAP

## 3. Anomaly Detection:

1. Isolation Forest
2. One-Class SVM (Support Vector Machine)
3. K-means clustering for anomaly detection


# PCA

- PCA is very simple in caret:

```
preprocess_method <- preProcess(  
  data, method = c("pca"))
```


```
data_pca <-  
predict(preprocess_method, data)
```

PCA Components for Pima Indians Diabetes Dataset



# PCA + Neural Network

We can do “automatic feature selection” by doing a PCA and then fit a ML model to the first  $k$  components:



# Practical Session 1.3

A first small project